

ML_Approaches_IoT-Real 2

May 13, 2024

```
[1]: %%time

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import scipy as sp
import seaborn as sns
import joblib
import os
from sklearn.feature_selection import f_classif, chi2, VarianceThreshold
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier, BaggingClassifier, \
    AdaBoostClassifier, StackingClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.impute import SimpleImputer
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix, f1_score, recall_score, precision_score

file_path = "/Users/ML-EdgeIIoT-dataset.csv"
data = pd.read_csv(file_path, low_memory=False)
```

CPU times: user 2.92 s, sys: 2.87 s, total: 5.79 s
Wall time: 3.19 s

```
[2]: %%time
# Columns to drop from DataFrame is named 'df'
columns_to_drop = [
    'frame.time', 'ip.src_host', 'ip.dst_host', 'arp.dst.proto_ipv4',
    'arp.src.proto_ipv4', 'icmp.checksum', 'icmp.seq_le', 'icmp.
    transmit_timestamp',
```

```

        'icmp.unused', 'http.file_data', 'http.request.uri.query', 'http.request.
        ↳method',
        'http.referer', 'http.request.full_uri', 'http.request.version', 'http.
        ↳response',
        'tcp.checksum', 'tcp.options', 'tcp.payload', 'udp.stream', 'udp.
        ↳time_delta',
        'dns.qry.name', 'dns.qry.name.len', 'dns.retransmission', 'dns.
        ↳retransmit_request',
        'dns.retransmit_request_in', 'mqtt.conack.flags', 'mqtt.msg_decoded_as',
        ↳'mqtt.msg',
        'mqtt.protoname', 'mqtt.topic', 'mbtcp.trans_id', 'mbtcp.unit_id'
    ]

# Drop the specified columns
df = data.drop(columns=columns_to_drop)

# Print the updated DataFrame
print(df.info())

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 157800 entries, 0 to 157799
```

```
Data columns (total 30 columns):
```

#	Column	Non-Null Count	Dtype
0	arp.opcode	157800 non-null	float64
1	arp.hw.size	157800 non-null	float64
2	http.content_length	157800 non-null	float64
3	http.tls_port	157800 non-null	float64
4	tcp.ack	157800 non-null	float64
5	tcp.ack_raw	157800 non-null	float64
6	tcp.connection.fin	157800 non-null	float64
7	tcp.connection.rst	157800 non-null	float64
8	tcp.connection.syn	157800 non-null	float64
9	tcp.connection.synack	157800 non-null	float64
10	tcp.dstport	157800 non-null	float64
11	tcp.flags	157800 non-null	float64
12	tcp.flags.ack	157800 non-null	float64
13	tcp.len	157800 non-null	float64
14	tcp.seq	157800 non-null	float64
15	tcp.srcport	157800 non-null	object
16	udp.port	157800 non-null	float64
17	dns.qry.qu	157800 non-null	float64
18	dns.qry.type	157800 non-null	float64
19	mqtt.conflag.cleansess	157800 non-null	float64
20	mqtt.conflags	157800 non-null	float64
21	mqtt.hdrflags	157800 non-null	float64
22	mqtt.len	157800 non-null	float64

```

23  mqtt.msgtype          157800 non-null  float64
24  mqtt.proto_len        157800 non-null  float64
25  mqtt.topic_len        157800 non-null  float64
26  mqtt.ver              157800 non-null  float64
27  mbtcp.len             157800 non-null  float64
28  Attack_label          157800 non-null  int64
29  Attack_type           157800 non-null  object
dtypes: float64(27), int64(1), object(2)
memory usage: 36.1+ MB
None
CPU times: user 24.1 ms, sys: 16.6 ms, total: 40.7 ms
Wall time: 56.5 ms

```

```

[3]: %%time
label_counts = df['Attack_type'].value_counts() #Total number of outut and
↪ frequencies
print(label_counts)

```

```

Attack_type
Normal          24301
DDoS_UDP        14498
DDoS_ICMP       14090
Ransomware      10925
DDoS_HTTP       10561
SQL_injection   10311
Uploading       10269
DDoS_TCP        10247
Backdoor        10195
Vulnerability_scanner 10076
Port_Scanning   10071
XSS             10052
Password        9989
MITM            1214
Fingerprinting  1001
Name: count, dtype: int64
CPU times: user 4.68 ms, sys: 238 µs, total: 4.92 ms
Wall time: 5.23 ms

```

```

[4]: %%time

# Define a regular expression pattern to match the entries to be eliminated
pattern = r'DESKTOP-UHFOSF2|DESKTOP-UHFOSF2\.local|_googlecast\.tcp\.local'

# Create a boolean mask indicating which rows have the specified entries in
↪ 'tcp.srcport'
mask = df['tcp.srcport'].str.contains(pattern, regex=True)

# Replace the specified entries with a default value (e.g., 0)

```

```

df.loc[mask, 'tcp.srcport'] = 0

# Convert the 'tcp.srcport' column to numeric data type
df['tcp.srcport'] = pd.to_numeric(df['tcp.srcport'], errors='coerce')

# Print the updated DataFrame
#print(df)

```

CPU times: user 59.6 ms, sys: 3.35 ms, total: 62.9 ms

Wall time: 71.3 ms

```

[5]: %%time
# Data features
X = df.drop(['Attack_label', 'Attack_type'], axis=1)

# Target Variable
y = df['Attack_type'] # for multi-class classification (types of attacks).

#X.info()

```

CPU times: user 4.28 ms, sys: 4.24 ms, total: 8.52 ms

Wall time: 8.02 ms

```

[6]: %%time

# Convert X to a DataFrame
#X = pd.DataFrame(X)

# Store the feature names
feature_names = X.columns

# Remove constant features
constant_filter = VarianceThreshold(threshold=0)
X = constant_filter.fit_transform(X)

# Perform ANOVA F-test
with np.errstate(divide='ignore', invalid='ignore'):
    f_scores, p_values = f_classif(X, y)

# Create a DataFrame to store the ANOVA F-test results
anova_results = pd.DataFrame({'Feature': feature_names[constant_filter.
    ↪get_support()], 'F-score': f_scores, 'p-value': p_values})

# Sort the results by F-score in descending order
anova_results = anova_results.sort_values('F-score', ascending=False)

print("ANOVA F-test results:")
print(anova_results)

```

```

# Perform Chi-Squared test
chi2_scores, p_values = chi2(X, y)

# Create a DataFrame to store the Chi-Squared test results
chi2_results = pd.DataFrame({'Feature': feature_names[constant_filter.
    ↳get_support()], 'Chi2-score': chi2_scores, 'p-value': p_values})

# Sort the results by Chi2-score in descending order
chi2_results = chi2_results.sort_values('Chi2-score', ascending=False)

print("\nChi-Squared test results:")
print(chi2_results)

```

ANOVA F-test results:

	Feature	F-score	p-value
10	tcp.flags	13121.586588	0.000000e+00
11	tcp.flags.ack	12038.253912	0.000000e+00
3	tcp.ack	6334.830157	0.000000e+00
4	tcp.ack_raw	5376.709017	0.000000e+00
7	tcp.connection.syn	3717.469027	0.000000e+00
6	tcp.connection.rst	3253.471612	0.000000e+00
9	tcp.dstport	3117.015249	0.000000e+00
14	tcp.srcport	2200.444757	0.000000e+00
16	dns.qry.qu	1919.975688	0.000000e+00
12	tcp.len	1358.312186	0.000000e+00
21	mqtt.msgtype	1053.240720	0.000000e+00
19	mqtt.hdrflags	1053.240720	0.000000e+00
5	tcp.connection.fin	979.716773	0.000000e+00
20	mqtt.len	904.167097	0.000000e+00
13	tcp.seq	883.056044	0.000000e+00
1	arp.hw.size	572.120040	0.000000e+00
18	mqtt.conflags	517.045704	0.000000e+00
22	mqtt.proto_len	517.045704	0.000000e+00
24	mqtt.ver	517.045704	0.000000e+00
17	mqtt.conflag.cleansess	517.045704	0.000000e+00
23	mqtt.topic_len	515.301739	0.000000e+00
0	arp.opcode	500.857426	0.000000e+00
8	tcp.connection.synack	450.238156	0.000000e+00
2	http.content_length	291.584976	0.000000e+00
15	udp.port	9.752936	3.536131e-22

Chi-Squared test results:

	Feature	Chi2-score	p-value
3	tcp.ack	7.626958e+13	0.0
4	tcp.ack_raw	6.297417e+13	0.0
13	tcp.seq	1.525879e+12	0.0
9	tcp.dstport	1.110271e+09	0.0

14	tcp.srcport	6.920185e+08	0.0
12	tcp.len	2.234175e+08	0.0
16	dns.qry.qu	1.569056e+07	0.0
2	http.content_length	1.426371e+07	0.0
15	udp.port	6.626143e+06	0.0
19	mqtt.hdrflags	2.237681e+06	0.0
10	tcp.flags	5.844437e+05	0.0
20	mqtt.len	3.635198e+05	0.0
23	mqtt.topic_len	1.642794e+05	0.0
21	mqtt.msgtype	1.398550e+05	0.0
1	arp.hw.size	4.528446e+04	0.0
7	tcp.connection.syn	3.413548e+04	0.0
6	tcp.connection.rst	3.202167e+04	0.0
11	tcp.flags.ack	2.972712e+04	0.0
22	mqtt.proto_len	2.746780e+04	0.0
24	mqtt.ver	2.746780e+04	0.0
18	mqtt.conflags	1.373390e+04	0.0
5	tcp.connection.fin	1.188649e+04	0.0
0	arp.opcode	1.061158e+04	0.0
17	mqtt.conflag.cleansess	6.866950e+03	0.0
8	tcp.connection.synack	5.880227e+03	0.0


CPU times: user 463 ms, sys: 42.3 ms, total: 506 ms
Wall time: 515 ms

```
[7]: %%time

scaler = StandardScaler()
X = scaler.fit_transform(X)

# Convert the NumPy array back to a DataFrame
df = pd.DataFrame(X)

# Visualize distribution of target variable
plt.figure(figsize=(8,6))
sns.countplot(y=y)
plt.title("Distribution of Attack Types")
plt.show()
```

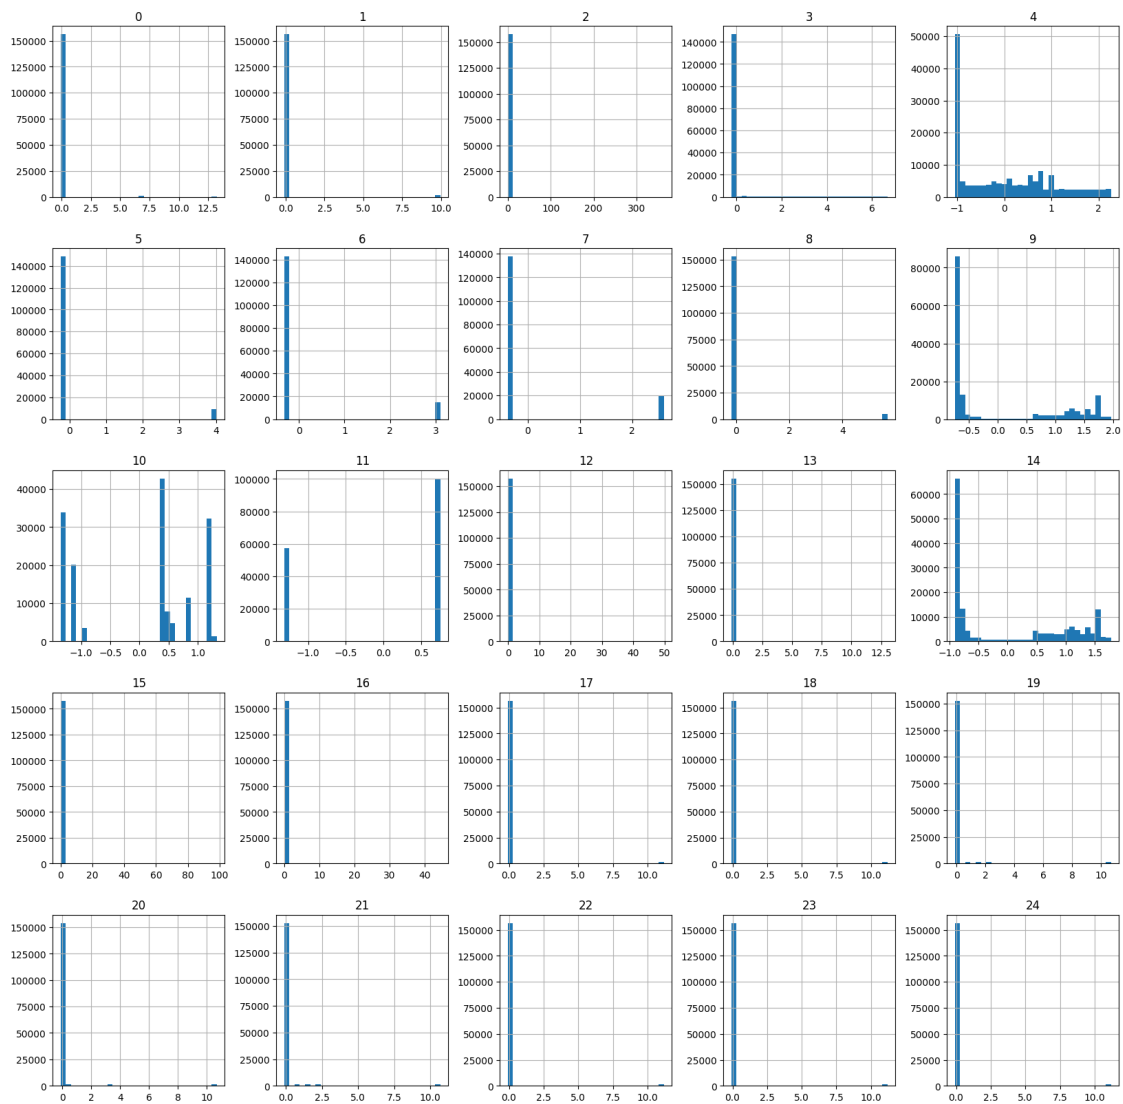


CPU times: user 706 ms, sys: 502 ms, total: 1.21 s

Wall time: 410 ms


```
[8]: # Plot histograms of features
df.hist(bins=30, figsize=(20,20))
plt.suptitle("Histograms of Features", fontsize=20)
plt.show()
```

Histograms of Features



```
[9]: # Compute correlation matrix
corr = df.corr()

# Plot correlation heatmap
plt.figure(figsize=(28,28))
sns.heatmap(corr, cmap='RdBu_r', annot=True, fmt=".2f", annot_kws={'size':10},
            xticklabels=corr.columns, yticklabels=corr.columns, square=True,
            linewidth=0.5)
plt.title("Correlation Heatmap of Features", fontsize=20)
plt.show()
```

[10]: df

[10]:

	0	1	2	3	4	5	6	\
0	-0.094772	-0.100375	-0.064076	-0.230878	-1.048497	-0.24846	-0.322332	
1	-0.094772	-0.100375	-0.064076	-0.230878	-1.048497	-0.24846	-0.322332	
2	-0.094772	-0.100375	-0.064076	-0.230878	-1.048497	-0.24846	-0.322332	
3	-0.094772	-0.100375	-0.064076	-0.230878	-1.048497	-0.24846	-0.322332	
4	-0.094772	-0.100375	-0.064076	-0.230878	-1.048497	-0.24846	-0.322332	
...	
157795	-0.094772	-0.100375	-0.064076	-0.230878	-1.048497	-0.24846	-0.322332	
157796	-0.094772	-0.100375	-0.064076	-0.230878	-1.048497	-0.24846	-0.322332	

```

157797 -0.094772 -0.100375 -0.064076 -0.230878 -1.048497 -0.24846 -0.322332
157798 -0.094772 -0.100375 -0.064076 -0.230878 -1.048497 -0.24846 -0.322332
157799 -0.094772 -0.100375 -0.064076 -0.230878 -1.048497 -0.24846 -0.322332

```

```

          7          8          9  ...          15          16          17  \
0      -0.382876 -0.17571 -0.74375 ... -0.012631 -0.033762 -0.089357
1      -0.382876 -0.17571 -0.74375 ... -0.012631 -0.033762 -0.089357
2      -0.382876 -0.17571 -0.74375 ... -0.012631 -0.033762 -0.089357
3      -0.382876 -0.17571 -0.74375 ... -0.012631 -0.033762 -0.089357
4      -0.382876 -0.17571 -0.74375 ... -0.012631 -0.033762 -0.089357
...
157795 -0.382876 -0.17571 -0.74375 ... -0.012631 -0.033762 -0.089357
157796 -0.382876 -0.17571 -0.74375 ... -0.012631 -0.033762 -0.089357
157797 -0.382876 -0.17571 -0.74375 ... -0.012631 -0.033762 -0.089357
157798 -0.382876 -0.17571 -0.74375 ... -0.012631 -0.033762 -0.089357
157799 -0.382876 -0.17571 -0.74375 ... -0.012631 -0.033762 -0.089357

          18          19          20          21          22          23          24
0      -0.089357 -0.124729 -0.116271 -0.124729 -0.089357 -0.089213 -0.089357
1      -0.089357 -0.124729 -0.116271 -0.124729 -0.089357 -0.089213 -0.089357
2      -0.089357 -0.124729 -0.116271 -0.124729 -0.089357 -0.089213 -0.089357
3      -0.089357 -0.124729 -0.116271 -0.124729 -0.089357 -0.089213 -0.089357
4      -0.089357 -0.124729 -0.116271 -0.124729 -0.089357 -0.089213 -0.089357
...
157795 -0.089357 -0.124729 -0.116271 -0.124729 -0.089357 -0.089213 -0.089357
157796 -0.089357 -0.124729 -0.116271 -0.124729 -0.089357 -0.089213 -0.089357
157797 -0.089357 -0.124729 -0.116271 -0.124729 -0.089357 -0.089213 -0.089357
157798 -0.089357 -0.124729 -0.116271 -0.124729 -0.089357 -0.089213 -0.089357
157799 -0.089357 -0.124729 -0.116271 -0.124729 -0.089357 -0.089213 -0.089357

```

[157800 rows x 25 columns]

```

[11]: %%time
# Split the data into training+validation and testing sets
X_train_val, X_test, y_train_val, y_test = train_test_split(X, y, test_size=0.
    ↪2, random_state=42)

# Split the training+validation set into training and validation sets
val_size = 0.2 # Validation set size (20% of the training+validation set)
X_train, X_val, y_train, y_val = train_test_split(X_train_val, y_train_val,
    ↪test_size=val_size, random_state=42)

# Print the sizes of the testing, validation, and training sets
print(f"Testing set size: {len(X_test)}")
print(f"Validation set size: {len(X_val)}")
print(f"Training set size: {len(X_train)}")

```

```

# Create an imputer object
imputer = SimpleImputer(strategy='mean') # or 'median', 'most_frequent'

# Impute missing values in the training, validation, and testing data
X_train_imputed = imputer.fit_transform(X_train)
X_val_imputed = imputer.transform(X_val)
X_test_imputed = imputer.transform(X_test)

# Create a Random Forest classifier
rf = RandomForestClassifier(n_estimators=100, random_state=42)

# Check if the trained Random Forest model exists
if os.path.exists('rf_model.joblib'):
    # Load the trained model from disk
    rf = joblib.load('rf_model.joblib')
else:
    # Fit the Random Forest model on the imputed training data
    rf.fit(X_train_imputed, y_train)
    # Save the trained model to disk
    joblib.dump(rf, 'rf_model.joblib')

# Get feature importances
importances = rf.feature_importances_

# Sort feature importances in descending order
indices = np.argsort(importances)[::-1]

# Get feature names
feature_names = df.columns

# Create a DataFrame with feature importances
feature_importances_df = pd.DataFrame({'Feature': feature_names, 'Importance':
    ↪ importances})

# Sort the DataFrame by importance in descending order
feature_importances_df = feature_importances_df.sort_values('Importance',
    ↪ ascending=False)

# Select the top N features (e.g., top 18)
top_n_features = 18
selected_features = feature_importances_df.head(top_n_features)

# Filter the imputed training, validation, and testing data with selected
    ↪ features
X_train_selected = X_train_imputed[:, selected_features.index]
X_val_selected = X_val_imputed[:, selected_features.index]
X_test_selected = X_test_imputed[:, selected_features.index]

```

```

# Initialize a DataFrame to store the accuracies
accuracies_df = pd.DataFrame(columns=['Model', 'Train Accuracy', 'Validation_
↳Accuracy', 'Test Accuracy'])

# SVM
if os.path.exists('svm_model.joblib'):
    svm = joblib.load('svm_model.joblib')
else:
    svm = SVC()
    svm.fit(X_train_selected, y_train)
    joblib.dump(svm, 'svm_model.joblib')

svm_score_train = svm.score(X_train_selected, y_train)
svm_score_val = svm.score(X_val_selected, y_val)
svm_score_test = svm.score(X_test_selected, y_test)
accuracies_df = pd.concat([accuracies_df, pd.DataFrame({'Model': ['SVM'],
↳'Train Accuracy': [svm_score_train], 'Validation Accuracy': [svm_score_val],
↳'Test Accuracy': [svm_score_test]})], ignore_index=True)

# KNN
if os.path.exists('knn_model.joblib'):
    knn = joblib.load('knn_model.joblib')
else:
    knn = KNeighborsClassifier()
    knn.fit(X_train_selected, y_train)
    joblib.dump(knn, 'knn_model.joblib')

knn_score_train = knn.score(X_train_selected, y_train)
knn_score_val = knn.score(X_val_selected, y_val)
knn_score_test = knn.score(X_test_selected, y_test)
accuracies_df = pd.concat([accuracies_df, pd.DataFrame({'Model': ['KNN'],
↳'Train Accuracy': [knn_score_train], 'Validation Accuracy': [knn_score_val],
↳'Test Accuracy': [knn_score_test]})], ignore_index=True)

# MLP
if os.path.exists('mlp_model.joblib'):
    mlp = joblib.load('mlp_model.joblib')
else:
    mlp = MLPClassifier(max_iter=500)
    mlp.fit(X_train_selected, y_train)
    joblib.dump(mlp, 'mlp_model.joblib')

mlp_score_train = mlp.score(X_train_selected, y_train)
mlp_score_val = mlp.score(X_val_selected, y_val)
mlp_score_test = mlp.score(X_test_selected, y_test)

```

```

accuracies_df = pd.concat([accuracies_df, pd.DataFrame({'Model': ['MLP'],
↳ 'Train Accuracy': [mlp_score_train], 'Validation Accuracy': [mlp_score_val],
↳ 'Test Accuracy': [mlp_score_test]})], ignore_index=True)

# Print the accuracies table
from tabulate import tabulate
print(tabulate(accuracies_df, headers='keys', tablefmt='psql'))

```

Testing set size: 31560

Validation set size: 25248

Training set size: 100992

<timed exec>:73: FutureWarning: The behavior of DataFrame concatenation with empty or all-NA entries is deprecated. In a future version, this will no longer exclude empty or all-NA columns when determining the result dtypes. To retain the old behavior, exclude the relevant entries before the concat operation.

	Model	Train Accuracy	Validation Accuracy	Test Accuracy
0	SVM	0.75304	0.752218	0.753169
1	KNN	0.881268	0.843988	0.844835
2	MLP	0.860167	0.857771	0.856305

CPU times: user 10min 56s, sys: 5.52 s, total: 11min 2s


Wall time: 9min 35s

```

[12]: from tabulate import tabulate

# Visualize the top 18 feature importances
plt.figure(figsize=(10, 8))
plt.title("Top 18 Feature Importances")
plt.barh(range(len(selected_features)), selected_features['Importance'],
↳ color='skyblue', align='center')
plt.yticks(range(len(selected_features)), selected_features['Feature'])
plt.xlabel("Importance")
plt.tight_layout()
plt.show()

```



```
[13]: %%time

## Split the data into training+validation and testing sets
X_train_val, X_test, y_train_val, y_test = train_test_split(X, y, test_size=0.
    ↪ 2, random_state=42)

# Create an imputer object
imputer = SimpleImputer(strategy='mean') # or 'median', 'most_frequent'

# Impute missing values in the training+validation and testing data
X_train_val_imputed = imputer.fit_transform(X_train_val)
X_test_imputed = imputer.transform(X_test)

# Split the imputed training+validation data into separate training and
    ↪ validation sets
X_train_imputed, X_val_imputed, y_train, y_val =
    ↪ train_test_split(X_train_val_imputed, y_train_val, test_size=0.2,
    ↪ random_state=42)
```

```

# Create a Random Forest classifier
rf = RandomForestClassifier(n_estimators=100, random_state=42, n_jobs=-1)

# Fit the Random Forest model on the imputed training+validation data
rf.fit(X_train_val_imputed, y_train_val)

# Get feature importances
importances = rf.feature_importances_

# Sort feature importances in descending order
indices = np.argsort(importances)[::-1]

# Get feature names
feature_names = df.columns

# Create a DataFrame with feature importances
feature_importances_df = pd.DataFrame({'Feature': feature_names, 'Importance': importances})

# Sort the DataFrame by importance in descending order
feature_importances_df.sort_values('Importance', ascending=False, inplace=True)

# Select the top N features (e.g., top 18)
top_n_features = 18
selected_features = feature_importances_df.head(top_n_features).index

# Filter the imputed training+validation and testing data with selected features
X_train_val_selected = X_train_val_imputed[:, selected_features]
X_test_selected = X_test_imputed[:, selected_features]

# Define parameter distributions for hyperparameter tuning
svm_param_dist = {'C': [0.1, 1, 10], 'kernel': ['linear', 'rbf']}
knn_param_dist = {'n_neighbors': [3, 5, 7], 'weights': ['uniform', 'distance']}
mlp_param_dist = {'hidden_layer_sizes': [(50,), (100,)], 'alpha': [0.0001, 0.001, 0.01]}

# Define base models
base_models = [
    ('SVM', SVC(probability=True), svm_param_dist),
    ('KNN', KNeighborsClassifier(), knn_param_dist),
    ('MLP', MLPClassifier(max_iter=500), mlp_param_dist)
]

def train_and_save_model(name, model, param_dist):
    print(f"Randomized search and hyperparameter tuning for {name}:")

```

```

    random_search = RandomizedSearchCV(estimator=model,
    ↪param_distributions=param_dist, n_iter=10, cv=5, random_state=42, n_jobs=-1)
    random_search.fit(X_train_val_selected, y_train_val)
    best_model = random_search.best_estimator_
    joblib.dump(best_model, f"{name.lower()}_model.joblib")
    print(f"Saved {name} model to file.")
    return best_model

tuned_base_models = [(name, joblib.load(f"{name.lower()}_model.joblib") if os.
    ↪path.exists(f"{name.lower()}_model.joblib") else train_and_save_model(name,
    ↪model, param_dist)) for name, model, param_dist in base_models]

# Create ensemble models
bagging = BaggingClassifier(estimator=SVC(probability=False), n_estimators=10,
    ↪random_state=42, n_jobs=-1)
boosting = AdaBoostClassifier(estimator=DecisionTreeClassifier(),
    ↪n_estimators=50, random_state=42)
stacking = StackingClassifier(estimators=tuned_base_models,
    ↪final_estimator=LogisticRegression(), cv=5, n_jobs=-1)
random_forest = RandomForestClassifier(n_estimators=100, random_state=42,
    ↪n_jobs=-1)
gradient_boosting = GradientBoostingClassifier(n_estimators=100,
    ↪random_state=42)

ensemble_models = [
    ('Bagging', bagging),
    ('Boosting', boosting),
    ('Random Forest', random_forest),
    ('Gradient Boosting', gradient_boosting),
    ('Stacking', stacking)
]

model_accuracies = {}

#def evaluate_models(models, X_val, y_val, X_test, y_test):
def evaluate_models(models, X_train, y_train, X_val, y_val, X_test, y_test):
    for name, model in models:
        model.fit(X_train_selected, y_train)

        # Evaluate on training set
        y_train_pred = model.predict(X_train)
        train_accuracy = accuracy_score(y_train, y_train_pred)
        print(f"{name} - Training accuracy: {train_accuracy:.3f}")

        # Evaluate on validation set
        y_val_pred = model.predict(X_val)

```



```

val_accuracy = accuracy_score(y_val, y_val_pred)
print(f"{name} - Validation accuracy: {val_accuracy:.3f}")

# Evaluate on test set
y_test_pred = model.predict(X_test)
test_accuracy = accuracy_score(y_test, y_test_pred)
model_accuracies[name] = test_accuracy
print(f"{name} - Test accuracy: {test_accuracy:.3f}")

# Compute the confusion matrix for test set
cm = confusion_matrix(y_test, y_test_pred)

# Get the unique class labels
class_labels = np.unique(y_test)

# Create a DataFrame from the confusion matrix
cm_df = pd.DataFrame(cm, index=class_labels, columns=class_labels)

# Create a figure and axes for the heatmap
fig, ax = plt.subplots(figsize=(10, 8))

# Create a heatmap using Seaborn
sns.heatmap(cm_df, annot=True, fmt='d', cmap='Blues', ax=ax, cbar=False)

# Set labels and title
ax.set_xlabel('Predicted')
ax.set_ylabel('Actual')
ax.set_title(f'Confusion Matrix - {name}')

# Display the heatmap
plt.show()

print("Classification Report:")
print(classification_report(y_test, y_test_pred))

# Calculate additional evaluation metrics
f1 = f1_score(y_test, y_test_pred, average='weighted') # or 'macro'
fnr = 1 - recall_score(y_test, y_test_pred, average='weighted') # or
↪ 'macro'
tpr = recall_score(y_test, y_test_pred, average='weighted') # or
↪ 'macro'
ppv = precision_score(y_test, y_test_pred, average='weighted') # or
↪ 'macro'

print(f"F1-score: {f1:.3f}")
print(f"FNR: {fnr:.3f}")
print(f"TPR: {tpr:.3f}")

```

```
print(f"PPV: {ppv:.3f}")
print()
```

... (rest of the code remains the same)

```
evaluate_models(tuned_base_models, X_train_imputed[:, selected_features],
↳ y_train, X_val_imputed[:, selected_features], y_val, X_test_selected, y_test)
evaluate_models(ensemble_models, X_train_imputed[:, selected_features],
↳ y_train, X_val_imputed[:, selected_features], y_val, X_test_selected, y_test)
```

SVM - Training accuracy: 0.747

SVM - Validation accuracy: 0.747

SVM - Test accuracy: 0.749

Actual \ Predicted	Backdoor	DDoS_HTTP	DDoS_ICMP	DDoS_TCP	DDoS_UDP	Fingerprinting	MITM	Normal	Password	Port_Scanning	Ransomware	SQL_injection	Uploading	Vulnerability_scanner	XSS
	Backdoor	DDoS_HTTP	DDoS_ICMP	DDoS_TCP	DDoS_UDP	Fingerprinting	MITM	Normal	Password	Port_Scanning	Ransomware	SQL_injection	Uploading	Vulnerability_scanner	XSS
Backdoor	1764	0	35	0	0	0	0	61	0	22	0	13	78	0	0
DDoS_HTTP	0	1215	0	0	0	0	0	323	12	0	0	450	6	34	89
DDoS_ICMP	0	0	2836	0	0	0	0	0	0	0	0	0	0	0	0
DDoS_TCP	0	20	0	1603	0	0	0	6	0	332	0	0	0	0	0
DDoS_UDP	0	0	384	0	2520	0	0	0	0	0	0	0	0	0	0
Fingerprinting	0	0	116	0	0	59	0	0	0	22	0	0	0	0	0
MITM	0	0	176	0	0	0	79	0	0	0	0	0	0	0	0
Normal	138	63	17	0	0	0	0	4609	0	10	0	0	140	0	8
Password	0	408	0	0	0	0	0	319	344	0	0	752	12	19	119
Port_Scanning	0	0	83	0	0	0	0	0	0	2045	0	0	0	0	0
Ransomware	0	0	133	0	0	0	0	0	0	113	1591	0	224	0	0
SQL_injection	0	246	2	0	0	0	0	0	0	0	0	1841	0	0	0
Uploading	7	163	3	0	0	0	0	325	20	0	7	503	662	1	324
Vulnerability_scanner	0	232	2	0	0	0	0	7	138	1	0	33	6	1590	0
XSS	0	217	87	0	0	0	0	644	3	10	0	206	12	0	866

Classification Report:

	precision	recall	f1-score	support
Backdoor	0.92	0.89	0.91	1973

DDoS_HTTP	0.47	0.57	0.52	2129
DDoS_ICMP	0.73	1.00	0.85	2836
DDoS_TCP	1.00	0.82	0.90	1961
DDoS_UDP	1.00	0.87	0.93	2904
Fingerprinting	1.00	0.30	0.46	197
MITM	1.00	0.31	0.47	255
Normal	0.73	0.92	0.82	4985
Password	0.67	0.17	0.28	1973
Port_Scanning	0.80	0.96	0.87	2128
Ransomware	1.00	0.77	0.87	2061
SQL_injection	0.48	0.88	0.63	2089
Uploading	0.58	0.33	0.42	2015
Vulnerability_scanner	0.97	0.79	0.87	2009
XSS	0.62	0.42	0.50	2045
accuracy			0.75	31560
macro avg	0.80	0.67	0.69	31560
weighted avg	0.77	0.75	0.73	31560

F1-score: 0.734

FNR: 0.251

TPR: 0.749

PPV: 0.771

KNN - Training accuracy: 0.877

KNN - Validation accuracy: 0.843

KNN - Test accuracy: 0.844

		Confusion Matrix - KNN														
Actual	Backdoor	1915	0	35	1	0	0	0	0	0	22	0	0	0	0	
	DDoS_HTTP	0	1485	0	0	0	0	5	174	0	0	221	94	21	129	
	DDoS_ICMP	0	0	2836	0	0	0	0	0	0	0	0	0	0	0	
	DDoS_TCP	0	1	0	1809	0	0	0	1	0	150	0	0	0	0	
	DDoS_UDP	0	0	278	0	2626	0	0	0	0	0	0	0	0	0	
	Fingerprinting	0	0	116	0	0	59	0	0	0	0	22	0	0	0	
	MITM	0	0	174	0	0	0	81	0	0	0	0	0	0	0	
	Normal	0	0	17	0	0	0	0	4943	8	0	10	0	6	0	1
	Password	0	364	0	0	0	0	0	17	796	0	0	420	116	153	107
	Port_Scanning	0	0	83	15	0	0	0	0	0	1860	170	0	0	0	0
	Ransomware	2	0	133	0	0	0	0	0	0	0	1926	0	0	0	0
	SQL_injection	0	297	2	0	0	0	0	0	213	0	0	1577	0	0	0
	Uploading	1	129	3	0	0	0	0	2	49	0	5	0	1327	0	499
	Vulnerability_scanner	0	20	2	0	0	0	0	3	87	0	1	3	3	1889	1
	XSS	0	106	87	0	0	0	0	0	36	0	10	2	286	1	1517
	Backdoor	DDoS_HTTP	DDoS_ICMP	DDoS_TCP	DDoS_UDP	Fingerprinting	MITM	Normal	Password	Port_Scanning	Ransomware	SQL_injection	Uploading	Vulnerability_scanner	XSS	
		Predicted														

Classification Report:

	precision	recall	f1-score	support
Backdoor	1.00	0.97	0.98	1973
DDoS_HTTP	0.62	0.70	0.66	2129
DDoS_ICMP	0.75	1.00	0.86	2836
DDoS_TCP	0.99	0.92	0.96	1961
DDoS_UDP	1.00	0.90	0.95	2904
Fingerprinting	1.00	0.30	0.46	197
MITM	1.00	0.32	0.48	255
Normal	0.99	0.99	0.99	4985
Password	0.58	0.40	0.48	1973
Port_Scanning	0.93	0.87	0.90	2128
Ransomware	0.89	0.93	0.91	2061
SQL_injection	0.71	0.75	0.73	2089
Uploading	0.72	0.66	0.69	2015
Vulnerability_scanner	0.92	0.94	0.93	2009
XSS	0.67	0.74	0.71	2045

accuracy			0.84	31560
macro avg	0.85	0.76	0.78	31560
weighted avg	0.85	0.84	0.84	31560

F1-score: 0.841

FNR: 0.156

TPR: 0.844

PPV: 0.849

MLP - Training accuracy: 0.823

MLP - Validation accuracy: 0.822

MLP - Test accuracy: 0.822

		Confusion Matrix - MLP														
Actual	Backdoor	1916	0	35	0	0	0	0	0	22	0	0	0	0	0	
	DDoS_HTTP	0	1434	0	0	0	0	0	208	0	0	202	64	56	165	
	DDoS_ICMP	0	0	2836	0	0	0	0	0	0	0	0	0	0	0	
	DDoS_TCP	0	0	0	1841	0	0	0	0	120	0	0	0	0	0	
	DDoS_UDP	0	0	286	0	2618	0	0	0	0	0	0	0	0	0	
	Fingerprinting	0	0	116	0	0	59	0	0	0	22	0	0	0	0	
	MITM	0	0	174	0	0	0	81	0	0	0	0	0	0	0	
	Normal	0	13	15	0	2	0	0	4587	103	10	0	0	245	10	0
	Password	0	412	0	0	0	0	0	837	0	0	411	40	144	129	
	Port_Scanning	0	0	83	23	0	0	0	0	2022	0	0	0	0	0	
	Ransomware	0	0	133	0	0	0	0	0	113	1598	46	171	0	0	
	SQL_injection	0	400	2	0	0	0	0	60	0	0	1627	0	0	0	
	Uploading	0	144	3	0	0	0	0	43	181	0	3	0	984	3	654
	Vulnerability_scanner	0	21	2	0	0	0	0	0	107	1	0	1	2	1875	0
	XSS	0	65	87	0	0	0	0	0	7	10	0	4	232	3	1637
	Backdoor	DDoS_HTTP	DDoS_ICMP	DDoS_TCP	DDoS_UDP	Fingerprinting	MITM	Normal	Password	Port_Scanning	Ransomware	SQL_injection	Uploading	Vulnerability_scanner	XSS	
	Predicted															

Classification Report:

	precision	recall	f1-score	support
Backdoor	1.00	0.97	0.99	1973

DDoS_HTTP	0.58	0.67	0.62	2129
DDoS_ICMP	0.75	1.00	0.86	2836
DDoS_TCP	0.99	0.94	0.96	1961
DDoS_UDP	1.00	0.90	0.95	2904
Fingerprinting	1.00	0.30	0.46	197
MITM	1.00	0.32	0.48	255
Normal	0.99	0.92	0.95	4985
Password	0.56	0.42	0.48	1973
Port_Scanning	0.87	0.95	0.91	2128
Ransomware	1.00	0.78	0.87	2061
SQL_injection	0.71	0.78	0.74	2089
Uploading	0.57	0.49	0.52	2015
Vulnerability_scanner	0.90	0.93	0.91	2009
XSS	0.63	0.80	0.71	2045
accuracy			0.82	31560
macro avg	0.84	0.74	0.76	31560
weighted avg	0.83	0.82	0.82	31560

F1-score: 0.820

FNR: 0.178

TPR: 0.822

PPV: 0.833

Bagging - Training accuracy: 0.747

Bagging - Validation accuracy: 0.747

Bagging - Test accuracy: 0.749

Actual	Backdoor	1764	0	35	0	0	0	0	61	0	22	0	13	78	0	0
	DDoS_HTTP	0	1221	0	0	0	0	0	324	13	0	0	448	6	34	83
	DDoS_ICMP	0	0	2836	0	0	0	0	0	0	0	0	0	0	0	0
	DDoS_TCP	0	19	0	1600	0	0	0	7	0	335	0	0	0	0	0
	DDoS_UDP	0	0	385	0	2519	0	0	0	0	0	0	0	0	0	0
	Fingerprinting	0	0	116	0	0	59	0	0	0	22	0	0	0	0	0
	MITM	0	0	176	0	0	0	79	0	0	0	0	0	0	0	0
	Normal	137	62	17	0	0	0	0	4610	2	10	0	0	141	0	6
	Password	0	412	0	0	0	0	0	318	345	0	0	752	12	18	116
	Port_Scanning	0	0	83	0	0	0	0	0	0	2045	0	0	0	0	0
	Ransomware	0	0	133	0	0	0	0	0	0	113	1591	0	224	0	0
	SQL_injection	0	248	2	0	0	0	0	0	0	0	0	1839	0	0	0
	Uploading	6	163	3	0	0	0	0	326	18	0	5	503	667	1	323
	Vulnerability_scanner	0	232	2	0	0	0	0	8	137	1	0	33	6	1590	0
	XSS	0	217	87	0	0	0	0	651	3	10	0	206	12	0	859
		Backdoor	DDoS_HTTP	DDoS_ICMP	DDoS_TCP	DDoS_UDP	Fingerprinting	MITM	Normal	Password	Port_Scanning	Ransomware	SQL_injection	Uploading	Vulnerability_scanner	XSS
		Predicted														

Classification Report:

	precision	recall	f1-score	support
Backdoor	0.93	0.89	0.91	1973
DDoS_HTTP	0.47	0.57	0.52	2129
DDoS_ICMP	0.73	1.00	0.85	2836
DDoS_TCP	1.00	0.82	0.90	1961
DDoS_UDP	1.00	0.87	0.93	2904
Fingerprinting	1.00	0.30	0.46	197
MITM	1.00	0.31	0.47	255
Normal	0.73	0.92	0.82	4985
Password	0.67	0.17	0.28	1973
Port_Scanning	0.80	0.96	0.87	2128
Ransomware	1.00	0.77	0.87	2061
SQL_injection	0.48	0.88	0.63	2089
Uploading	0.58	0.33	0.42	2015
Vulnerability_scanner	0.97	0.79	0.87	2009
XSS	0.62	0.42	0.50	2045

accuracy			0.75	31560
macro avg	0.80	0.67	0.69	31560
weighted avg	0.77	0.75	0.73	31560

F1-score: 0.734

FNR: 0.251

TPR: 0.749

PPV: 0.771

/Users/ebimol/miniforge3/lib/python3.10/site-packages/sklearn/ensemble/_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
 warnings.warn(

Boosting - Training accuracy: 0.948

Boosting - Validation accuracy: 0.910

Boosting - Test accuracy: 0.909

		Confusion Matrix - Boosting														
Actual	Backdoor -	1916	0	35	0	0	0	0	0	0	22	0	0	0	0	0
	DDoS_HTTP -	0	1747	0	0	0	0	0	0	179	0	0	128	27	13	35
	DDoS_ICMP -	0	0	2836	0	0	0	0	0	0	0	0	0	0	0	0
	DDoS_TCP -	0	0	0	1961	0	0	0	0	0	0	0	0	0	0	0
	DDoS_UDP -	0	0	277	0	2627	0	0	0	0	0	0	0	0	0	0
	Fingerprinting -	2	0	116	0	0	57	0	0	0	22	0	0	0	0	0
	MITM -	0	0	174	0	0	0	81	0	0	0	0	0	0	0	0
	Normal -	0	0	23	0	0	0	0	4952	0	10	0	0	0	0	0
	Password -	0	114	0	0	0	0	0	0	1694	0	0	90	27	28	20
	Port_Scanning -	0	0	83	0	0	0	0	0	0	2045	0	0	0	0	0
	Ransomware -	0	0	133	0	0	0	0	0	0	113	1815	0	0	0	0
	SQL_injection -	0	164	2	0	0	0	0	0	241	0	0	1682	0	0	0
	Uploading -	0	32	3	0	0	0	0	0	113	0	0	0	1723	0	144
	Vulnerability_scanner -	0	10	2	0	0	0	0	0	37	1	0	0	0	1958	1
	XSS -	0	71	87	0	0	0	0	0	155	10	0	1	133	0	1588
		Backdoor -	DDoS_HTTP -	DDoS_ICMP -	DDoS_TCP -	DDoS_UDP -	Fingerprinting -	MITM -	Normal -	Password -	Port_Scanning -	Ransomware -	SQL_injection -	Uploading -	Vulnerability_scanner -	XSS -
		Predicted														

Classification Report:

	precision	recall	f1-score	support
Backdoor	1.00	0.97	0.98	1973
DDoS_HTTP	0.82	0.82	0.82	2129
DDoS_ICMP	0.75	1.00	0.86	2836
DDoS_TCP	1.00	1.00	1.00	1961
DDoS_UDP	1.00	0.90	0.95	2904
Fingerprinting	1.00	0.29	0.45	197
MITM	1.00	0.32	0.48	255
Normal	1.00	0.99	1.00	4985
Password	0.70	0.86	0.77	1973
Port_Scanning	0.92	0.96	0.94	2128
Ransomware	1.00	0.88	0.94	2061
SQL_injection	0.88	0.81	0.84	2089
Uploading	0.90	0.86	0.88	2015
Vulnerability_scanner	0.98	0.97	0.98	2009
XSS	0.89	0.78	0.83	2045
accuracy			0.91	31560
macro avg	0.92	0.83	0.85	31560
weighted avg	0.92	0.91	0.91	31560

F1-score: 0.908

FNR: 0.091

TPR: 0.909

PPV: 0.919

Random Forest - Training accuracy: 0.951

Random Forest - Validation accuracy: 0.914

Random Forest - Test accuracy: 0.914

		Confusion Matrix - Random Forest														
Actual	Backdoor	1916	0	35	0	0	0	0	0	22	0	0	0	0	0	
	DDoS_HTTP	0	1744	0	0	0	0	0	159	0	0	143	30	8	45	
	DDoS_ICMP	0	0	2836	0	0	0	0	0	0	0	0	0	0	0	
	DDoS_TCP	0	0	0	1961	0	0	0	0	0	0	0	0	0	0	
	DDoS_UDP	0	0	277	0	2627	0	0	0	0	0	0	0	0	0	
	Fingerprinting	0	0	116	0	0	59	0	0	0	22	0	0	0	0	
	MITM	0	0	174	0	0	0	81	0	0	0	0	0	0	0	
	Normal	0	0	23	0	0	0	0	4952	0	10	0	0	0	0	
	Password	0	113	0	0	0	0	0	1681	0	0	101	25	21	32	
	Port_Scanning	0	0	83	0	0	0	0	0	2045	0	0	0	0	0	
	Ransomware	0	0	133	0	0	0	0	0	113	1815	0	0	0	0	
	SQL_injection	0	167	2	0	0	0	0	229	0	0	1691	0	0	0	
	Uploading	0	24	3	0	0	0	0	44	0	0	0	1742	0	202	
	Vulnerability_scanner	0	12	2	0	0	0	0	38	1	0	0	1	1953	2	
	XSS	0	48	87	0	0	0	0	42	10	0	0	124	1	1733	
	Backdoor	DDoS_HTTP	DDoS_ICMP	DDoS_TCP	DDoS_UDP	Fingerprinting	MITM	Normal	Password	Port_Scanning	Ransomware	SQL_injection	Uploading	Vulnerability_scanner	XSS	
	Predicted															

Classification Report:

	precision	recall	f1-score	support
Backdoor	1.00	0.97	0.99	1973
DDoS_HTTP	0.83	0.82	0.82	2129
DDoS_ICMP	0.75	1.00	0.86	2836
DDoS_TCP	1.00	1.00	1.00	1961
DDoS_UDP	1.00	0.90	0.95	2904
Fingerprinting	1.00	0.30	0.46	197
MITM	1.00	0.32	0.48	255
Normal	1.00	0.99	1.00	4985
Password	0.77	0.85	0.81	1973
Port_Scanning	0.92	0.96	0.94	2128
Ransomware	1.00	0.88	0.94	2061
SQL_injection	0.87	0.81	0.84	2089
Uploading	0.91	0.86	0.88	2015
Vulnerability_scanner	0.98	0.97	0.98	2009
XSS	0.86	0.85	0.85	2045

accuracy			0.91	31560
macro avg	0.93	0.83	0.85	31560
weighted avg	0.92	0.91	0.91	31560

F1-score: 0.913

FNR: 0.086

TPR: 0.914

PPV: 0.922

Gradient Boosting - Training accuracy: 0.921

Gradient Boosting - Validation accuracy: 0.920

Gradient Boosting - Test accuracy: 0.920

		Confusion Matrix - Gradient Boosting														
Actual	Backdoor	1916	0	35	0	0	0	0	0	22	0	0	0	0	0	0
	DDoS_HTTP	0	1845	0	0	0	0	0	0	33	0	0	139	33	11	68
	DDoS_ICMP	0	0	2836	0	0	0	0	0	0	0	0	0	0	0	0
	DDoS_TCP	0	0	0	1961	0	0	0	0	0	0	0	0	0	0	0
	DDoS_UDP	0	0	277	0	2627	0	0	0	0	0	0	0	0	0	0
	Fingerprinting	0	0	116	0	0	59	0	0	0	22	0	0	0	0	0
	MITM	0	0	174	0	0	0	81	0	0	0	0	0	0	0	0
	Normal	0	0	23	0	0	0	0	4952	0	10	0	0	0	0	0
	Password	0	217	0	0	0	0	0	0	1499	0	0	169	31	10	47
	Port_Scanning	0	0	83	0	0	0	0	0	0	2044	0	0	1	0	0
	Ransomware	0	0	133	0	0	0	0	0	0	113	1815	0	0	0	0
	SQL_injection	0	223	2	0	0	0	0	0	2	0	0	1862	0	0	0
	Uploading	0	38	3	0	0	0	0	0	0	0	0	0	1680	0	294
	Vulnerability_scanner	0	14	2	0	0	0	0	0	51	1	0	2	0	1938	1
	XSS	0	5	87	0	0	0	0	0	3	10	0	1	3	2	1934
		Backdoor	DDoS_HTTP	DDoS_ICMP	DDoS_TCP	DDoS_UDP	Fingerprinting	MITM	Normal	Password	Port_Scanning	Ransomware	SQL_injection	Uploading	Vulnerability_scanner	XSS
		Predicted														

Classification Report:

	precision	recall	f1-score	support
Backdoor	1.00	0.97	0.99	1973

DDoS_HTTP	0.79	0.87	0.83	2129
DDoS_ICMP	0.75	1.00	0.86	2836
DDoS_TCP	1.00	1.00	1.00	1961
DDoS_UDP	1.00	0.90	0.95	2904
Fingerprinting	1.00	0.30	0.46	197
MITM	1.00	0.32	0.48	255
Normal	1.00	0.99	1.00	4985
Password	0.94	0.76	0.84	1973
Port_Scanning	0.92	0.96	0.94	2128
Ransomware	1.00	0.88	0.94	2061
SQL_injection	0.86	0.89	0.87	2089
Uploading	0.96	0.83	0.89	2015
Vulnerability_scanner	0.99	0.96	0.98	2009
XSS	0.83	0.95	0.88	2045
accuracy			0.92	31560
macro avg	0.94	0.84	0.86	31560
weighted avg	0.93	0.92	0.92	31560

F1-score: 0.919

FNR: 0.080

TPR: 0.920

PPV: 0.930

/Users/ebimol/miniforge3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Stacking - Training accuracy: 0.863

Stacking - Validation accuracy: 0.850

Stacking - Test accuracy: 0.848

Actual	Backdoor	1892	0	35	5	0	0	0	2	0	11	15	0	13	0	0
	DDoS_HTTP	0	1491	0	0	0	0	0	0	188	0	0	233	41	16	160
	DDoS_ICMP	0	0	2836	0	0	0	0	0	0	0	0	0	0	0	0
	DDoS_TCP	0	3	0	1824	0	0	0	1	8	120	0	3	0	0	2
	DDoS_UDP	0	0	278	0	2626	0	0	0	0	0	0	0	0	0	0
	Fingerprinting	0	0	116	0	0	59	0	0	0	10	12	0	0	0	0
	MITM	0	0	176	0	0	0	79	0	0	0	0	0	0	0	0
	Normal	0	1	17	0	0	0	0	4943	8	1	9	0	0	0	6
	Password	0	371	0	0	0	0	0	0	957	0	0	394	52	73	126
	Port_Scanning	0	0	83	40	0	0	0	0	0	1903	102	0	0	0	0
	Ransomware	0	0	133	0	0	0	0	0	0	52	1876	0	0	0	0
	SQL_injection	0	233	2	0	0	0	0	0	280	0	0	1574	0	0	0
	Uploading	0	160	3	0	0	0	0	17	58	0	5	0	1110	0	662
	Vulnerability_scanner	0	6	2	0	0	0	0	1	164	0	1	10	3	1822	0
	XSS	0	46	87	0	0	0	0	0	14	6	4	4	101	1	1782
		Backdoor	DDoS_HTTP	DDoS_ICMP	DDoS_TCP	DDoS_UDP	Fingerprinting	MITM	Normal	Password	Port_Scanning	Ransomware	SQL_injection	Uploading	Vulnerability_scanner	XSS
		Predicted														

Classification Report:

	precision	recall	f1-score	support
Backdoor	1.00	0.96	0.98	1973
DDoS_HTTP	0.65	0.70	0.67	2129
DDoS_ICMP	0.75	1.00	0.86	2836
DDoS_TCP	0.98	0.93	0.95	1961
DDoS_UDP	1.00	0.90	0.95	2904
Fingerprinting	1.00	0.30	0.46	197
MITM	1.00	0.31	0.47	255
Normal	1.00	0.99	0.99	4985
Password	0.57	0.49	0.52	1973
Port_Scanning	0.90	0.89	0.90	2128
Ransomware	0.93	0.91	0.92	2061
SQL_injection	0.71	0.75	0.73	2089
Uploading	0.84	0.55	0.67	2015
Vulnerability_scanner	0.95	0.91	0.93	2009
XSS	0.65	0.87	0.75	2045

accuracy			0.85	31560
macro avg	0.86	0.76	0.78	31560
weighted avg	0.86	0.85	0.85	31560

F1-score: 0.846

FNR: 0.152

TPR: 0.848

PPV: 0.859

CPU times: user 2h 20min 36s, sys: 6min 35s, total: 2h 27min 12s

Wall time: 1h 20min 45s


```
[14]: import matplotlib.colors as mcolors

# Define the list of colors
color_list = ['red', 'green', 'yellow', 'cyan', 'black', 'lightgreen', 'pink',
              ↪ 'blue']

plt.figure(figsize=(10, 6))
bars = plt.bar(model_accuracies.keys(), model_accuracies.values(),
              ↪ color=color_list)
plt.ylim(0, 1)
plt.title("Model Performance Comparison")
plt.xlabel("Model")
plt.ylabel("Accuracy")
plt.xticks(rotation=45)

# Add count values on top of the bars
for bar in bars:
    height = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2, height,
             ↪ f'{height:.2f}', ha='center', va='bottom')

plt.show()
```



[]: